



Published in Towards Data Science

You have **2** free member-only stories left this month.

[Sign up for Medium and get an extra one](#)



Eligijus Bujokas

[Follow](#)

Apr 14, 2021 · 8 min read · ✨ · [Listen](#)




Save



Decision Tree Algorithm in Python From Scratch

Coding the popular algorithm using just NumPy and Pandas in Python and explaining what's under the hood



Decision tree schema; graph by author

The aim of this article is to make all the parts of a decision tree classifier clear by walking through the code that implements the algorithm. The code uses only NumPy, Pandas and the standard python libraries.

The full code can be accessed via <https://github.com/Eligijus112/decision-tree-python>

As of now, the code creates a decision tree when the target variable is binary and the features are numeric. This is completely sufficient to understand the algorithm.

The golden standard of building decision trees in python is the scikit-learn implementation:

--	--


1.10. Decision Trees - scikit-learn 0.24.1 documentation

Decision Trees (DTs) are a non-parametric supervised learning method used for classification and regression . The goal...

scikit-learn.org


When I tested out my code I wanted to make sure that the results are identical to the scikit-learn implementation.

The data used in this article is the famous Titanic survivor dataset. We will use two numeric variables — Age of the passenger and the Fare of the ticket — to predicting whether a passenger survived or not.



Age + Fare ~ Survival; Graph by author

The goal is to create the “best” splits of the numeric variables. Just eyeballing the data, we could guess that one good split is to split the data into two parts: observations that have $\text{Age} < 10$ and observations that have $\text{Age} \geq 10$:



Splitting the dataset; Graph by author

Now some immediate questions may rise:

Is this a good split?

Maybe a split at the Fare = 200 is a better one?

How do we quantify the “goodness” of a split?

How does a computer search for the best split?

All of these questions will be answered by the end of this article.

A decision tree algorithm (DT for short) is a machine learning algorithm that is used in classifying an observation given a set of input features. The algorithm creates a set of rules at various decision levels such that a certain metric is optimized.

The target variable will be denoted as $Y = \{0, 1\}$ and the feature matrix will be denoted as X .

Keywords to expand on:

Node

Gini impurity (a metric which we are optimizing)

Level

Splitting

A **node** is the building block in the decision tree. When viewing a typical schema of a decision tree (like the one in the title picture) the nodes are the rectangles or bubbles that have a downward connection to other nodes.

Open in app ↗ [Sign up](#) [Sign In](#)



Number of observations

The number of observations belonging to each of the binary target classes.

The feature matrix X representing the observations that fall into the node.

The custom **Node** class in python (written by me):

```

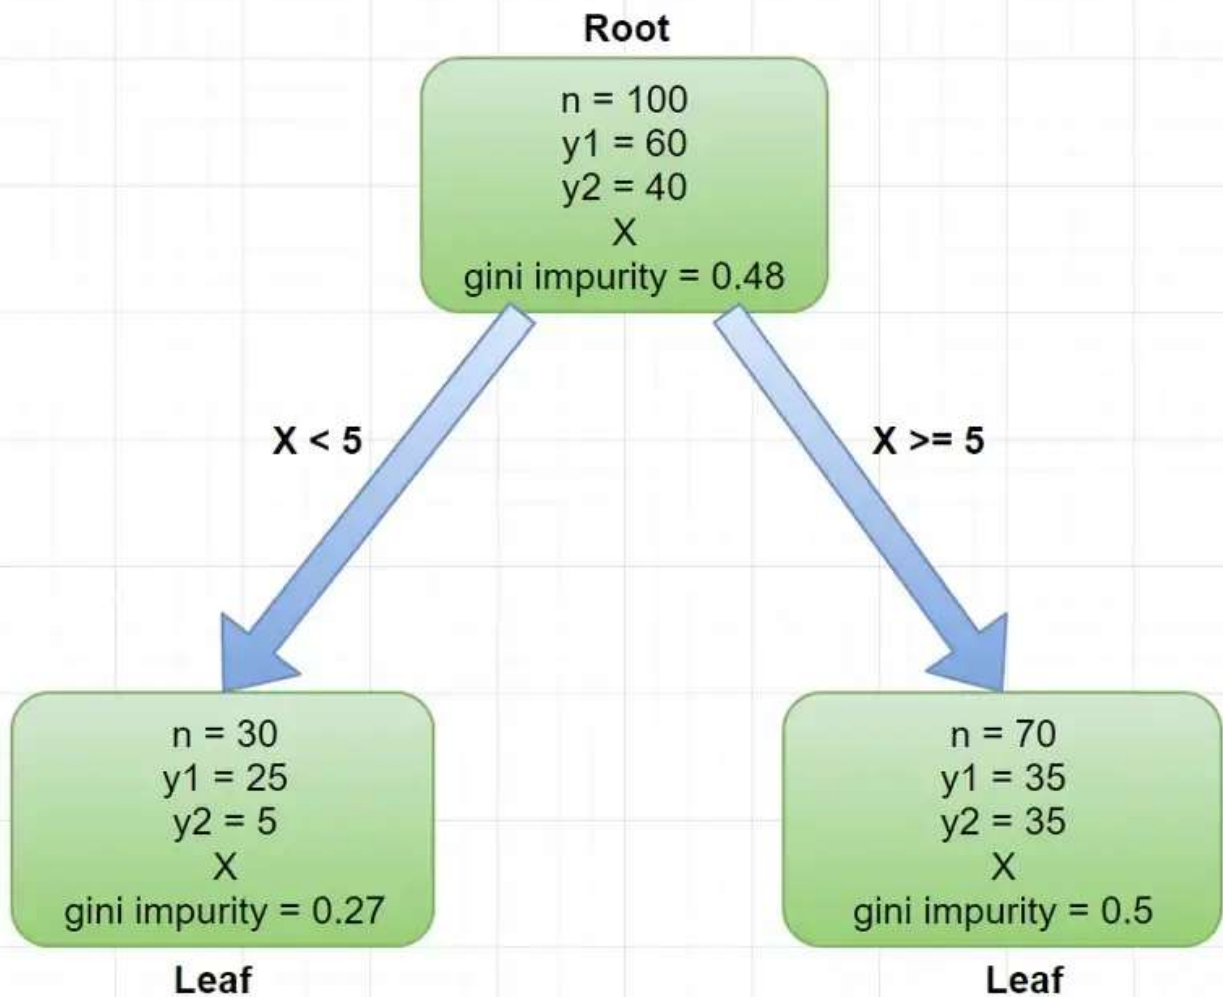
1  # Data wrangling
2  import pandas as pd
3
4  # Array math
5  import numpy as np
6
7  # Quick value count calculator
8  from collections import Counter
9
10
11 class Node:
12     """
13     Class for creating the nodes for a decision tree
14     """
15     def __init__(
16         self,
17         Y: list,
18         X: pd.DataFrame,
19         min_samples_split=None,
20         max_depth=None,
21         depth=None,
22         node_type=None,
23         rule=None
24     ):
25         # Saving the data to the node
26         self.Y = Y
27         self.X = X
28
29         # Saving the hyper parameters
30         self.min_samples_split = min_samples_split if min_samples_split else 20
31         self.max_depth = max_depth if max_depth else 5
32
33         # Default current depth of node
34         self.depth = depth if depth else 0
35
36         # Extracting all the features
37         self.features = list(self.X.columns)
38
39         # Type of node
40         self.node_type = node_type if node_type else 'root'
41
42         # Rule for splitting
43         self.rule = rule if rule else ""
44
45         # Calculating the counts of Y in the node
46         self.counts = Counter(Y)
47
48         # Getting the GINI impurity based on the Y distribution

```

```

48     # Getting the gini impurity based on the y distribution
49     self.gini_impurity = self.get_GINI()
50
51     # Sorting the counts and saving the final prediction of the node
52     counts_sorted = list(sorted(self.counts.items(), key=lambda item: item[1]))
53
54     # Getting the last item
55     yhat = None
56     if len(counts_sorted) > 0:

```



```

81
82     if y2_count is None:
83         y2_count = 0
84
85     # Getting the total observations
86     n = y1_count + y2_count
87
88     # If n is 0 then we return the lowest possible gini impurity
89     if n == 0:
90         return 0.0
91
92     # Getting the probability to see each of the classes
93     p1 = y1_count / n
94     p2 = y2_count / n
95

```

```

96         # Calculating GINI
97         gini = 1 - (p1 ** 2 + p2 ** 2)
98
99         # Returning the gini impurity
100         return gini
101
102     @staticmethod
103     def ma(x: np.array, window: int) -> np.array:
104         """
105         Calculates the moving average of the given list.
106         """
107         return np.convolve(x, np.ones(window), 'valid') / window
108
109     def get_GINI(self):
110         """
111         Function to calculate the GINI impurity of a node
112         """

```

Suppose we have two classes in the dataset:

$$k_1, k_2$$

Each of the classes have n_1 and n_2 observations.

The probability of observing something from one of the k classes is:

$$p(i) = P(x_i \in k_i) = \frac{n_i}{n_1 + n_2}, i \in \{1, 2\}$$

The GINI impurity of such a system is calculated with the following formula:

$$G = 1 - \sum_{i=1}^2 p(i)^2$$

```

126         df['Y'] = self.Y
127
128         # Getting the GINI impurity for the base input
129         GINI_base = self.get_GINI()
130
131         # Finding which split yields the best GINI gain
132         max_gain = 0
133
134         # Default best feature and split
135         best_feature = None
136         best_value = None
137
138         for feature in self.features:
139             # Dropping missing values
140             Xdf = df.dropna().sort_values(feature)
141
142             # Sorting the values and getting the rolling average
143             xmeans = self.ma(Xdf[feature].unique(), 2)

```



```

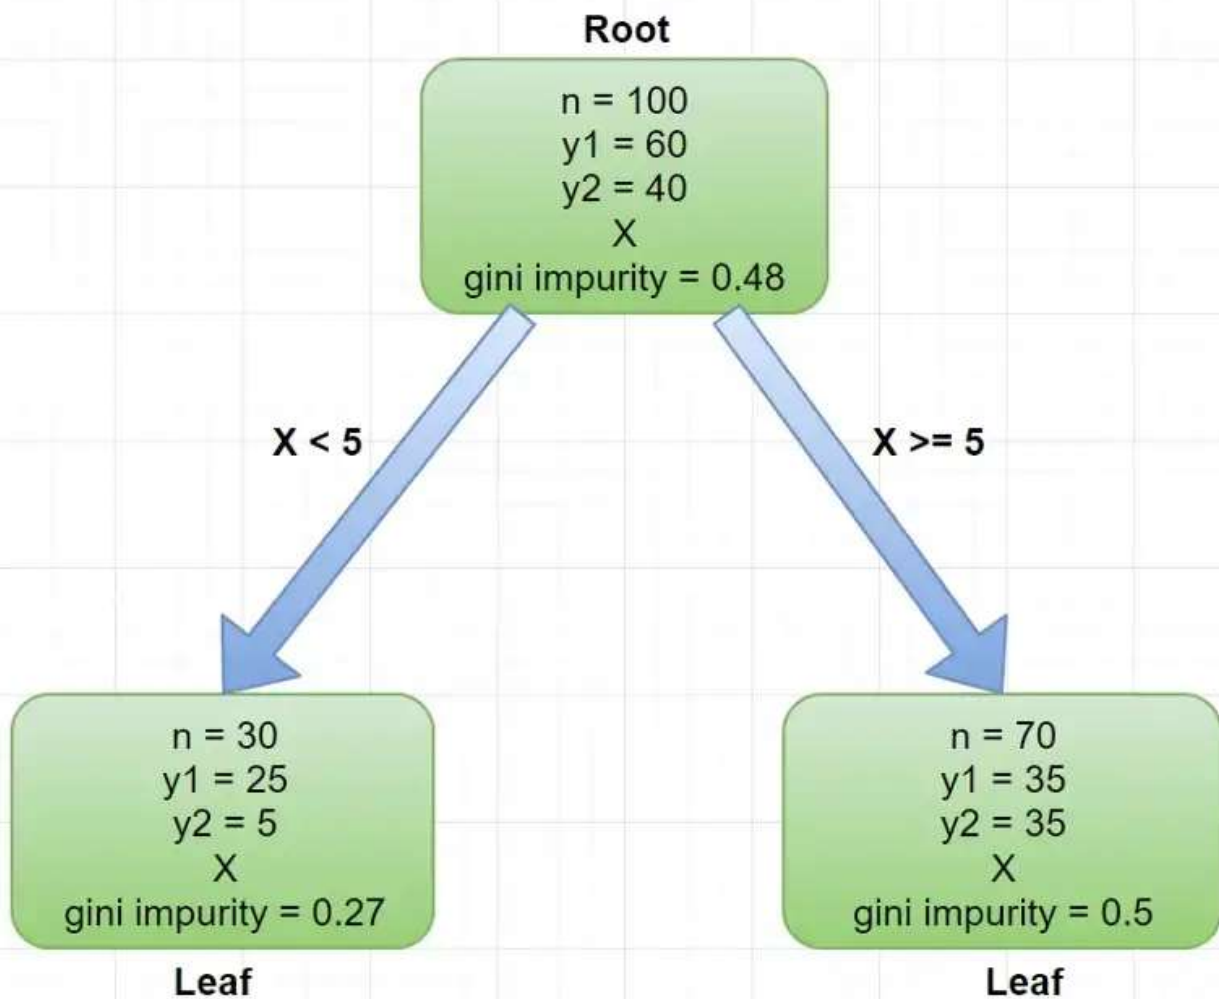
143         xmeans = self.find(Xdf[feature].unique(), 2,
144
145         for value in xmeans:
146             # Splitting the dataset
147             left_counts = Counter(Xdf[Xdf[feature]<value]['Y'])
148             right_counts = Counter(Xdf[Xdf[feature]>=value]['Y'])
149
150             # Getting the Y distribution from the dicts
151             y0_left, y1_left, y0_right, y1_right = left_counts.get(0, 0), left_counts.get(
152
153             # Getting the left and right gini impurities
154             gini_left = self.GINI_impurity(y0_left, y1_left)
155             gini_right = self.GINI_impurity(y0_right, y1_right)
156
157             # Getting the obs count from the left and the right data splits
158             n_left = y0_left + y1_left
159             n_right = y0_right + y1_right
160
161             # Calculating the weights for each of the nodes
162             w_left = n_left / (n_left + n_right)
163             w_right = n_right / (n_left + n_right)
164
165             # Calculating the weighted GINI impurity
166             wGINI = w_left * gini_left + w_right * gini_right
167
168             # Calculating the GINI gain
169             GINIGain = GINI_base - wGINI
170
171             # Checking if this is the best split so far
172             if GINIGain > max_gain:
173                 best_feature = feature
174                 best_value = value
175
176             # Setting the best gain to the current one
177             max_gain = GINIGain
178
179         return (best_feature, best_value)
180
181     def grow_tree(self):
182         """
183         Recursive method to create the decision tree
184         """
185         # Making a df from the data
186         df = self.X.copy()
187         df['Y'] = self.Y
188
189         # If there is GINI to be gained, we split further
190         if (self.depth < self.max_depth) and (self.n >= self.min_samples_split):

```

```

191
192         # Getting the best split
193         best_feature, best_value = self.best_split()
194

```



```

219         right_df[self.features],
220         depth=self.depth + 1,
221         max_depth=self.max_depth,
222         min_samples_split=self.min_samples_split,
223         node_type='right_node',
224         rule=f"{best_feature} > {round(best_value, 3)}"
225     )
226
227     self.right = right
228     self.right.grow_tree()
229
230     def print_info(self, width=4):
231         """

```

$$GINI_{gain} = \Delta Gini = Gini_{parent} - \left(Gini_{left} \frac{n_{left}}{n_{right} + n_{left}} + Gini_{right} \frac{n_{right}}{n_{right} + n_{left}} \right)$$

```

234         # Defining the number of spaces
235         const = int(self.depth * width ** 1.5)
236         spaces = "-" * const
237
238         if self.node type == 'root':

```

```

239         print("Root")
240     else:
241         print(f"|{spaces} Split rule: {self.rule}")
242     print(f"{' ' * const} | GINI impurity of the node: {round(self.gini_impurity, 2)}")
243     print(f"{' ' * const} | Class distribution in the node: {dict(self.counts)}")
244     print(f"{' ' * const} | Predicted class: {self.yhat}")
245
246     def print_tree(self):
247         """
248         Prints the whole tree from the current node to the bottom
249         """
250         self.print_info()
251
252         if self.left is not None:
253             self.left.print_tree()
254
255         if self.right is not None:
256             self.right.print_tree()
257
258     def predict(self, X:pd.DataFrame):
259         """
260         Batch prediction method
261         """
262         predictions = []
263
264         for _, x in X.iterrows():
265             values = {}
266             for feature in self.features:
267                 values.update({feature: x[feature]})
268
269             predictions.append(self.predict_obs(values))
270
271         return predictions
272
273     def predict_obs(self, values: dict) -> int:
274         """
275         Method to predict the class given a set of features
276         """
277         cur_node = self
278         while cur_node.depth < cur_node.max_depth:
279             # Traversing the nodes all the way to the bottom
280             best_feature = cur_node.best_feature
281             best_value = cur_node.best_value
282
283             if cur_node.n < cur_node.min_samples_split:
284                 break
285

```

```

286         if (values.get(best_feature) < best_value):
287             if self.left is not None:
288                 cur_node = cur_node.left
289             else:
290                 if self.right is not None:
291                     cur_node = cur_node.right

```

Here for the best split search to start to right example, if the node has 51 observations but the min_samples_split = 55, then the growth of the tree stops.

So, how does the code work?

First of all, read the data:

```

# Loading data
d = pd.read_csv('data/train.csv')

# Dropping missing values
dtree = d[['Survived', 'Age', 'Fare']].dropna().copy()

# Defining the X and Y matrices
Y = dtree['Survived'].values
X = dtree[['Age', 'Fare']]

# Saving the feature list
features = list(X.columns)

```

Then we define the dictionary of hyperparameters.

```

hp = {
    'max_depth': 3,
    'min_samples_split': 50
}

```

Then we initiate the root node:

```

root = Node(Y, X, **hp)

```

The main tree building function is the **grow_tree()** function.

```
root.grow_tree()
```

And that's it!

To view the results, we can invoke the **print_tree()** function.

```
root.print_tree()
```

The results:

```

Root
| GINI impurity of the node: 0.48
| Class distribution in the node: {0: 424, 1: 290}
| Predicted class: 0
|----- Split rule: Fare <= 52.277
|       | GINI impurity of the node: 0.44
|       | Class distribution in the node: {0: 389, 1: 195}
|       | Predicted class: 0
|----- Split rule: Fare <= 10.481
|       | GINI impurity of the node: 0.32
|       | Class distribution in the node: {0: 192, 1: 47}
|       | Predicted class: 0
|----- Split rule: Age <= 32.5
|       | GINI impurity of the node: 0.37
|       | Class distribution in the node: {0: 134, 1: 43}
|       | Predicted class: 0
|----- Split rule: Age > 32.5
|       | GINI impurity of the node: 0.12
|       | Class distribution in the node: {0: 58, 1: 4}
|       | Predicted class: 0
|----- Split rule: Fare > 10.481
|       | GINI impurity of the node: 0.49
|       | Class distribution in the node: {0: 197, 1: 148}
|       | Predicted class: 0
|----- Split rule: Age <= 6.5
|       | GINI impurity of the node: 0.41
|       | Class distribution in the node: {0: 12, 1: 30}
|       | Predicted class: 1
|----- Split rule: Age > 6.5
|       | GINI impurity of the node: 0.48
|       | Class distribution in the node: {0: 185, 1: 118}
|       | Predicted class: 0
|----- Split rule: Fare > 52.277
|       | GINI impurity of the node: 0.39
|       | Class distribution in the node: {1: 95, 0: 35}
|       | Predicted class: 1
|----- Split rule: Age <= 63.5
|       | GINI impurity of the node: 0.38
|       | Class distribution in the node: {1: 95, 0: 32}
|       | Predicted class: 1
|----- Split rule: Age <= 29.5
|       | GINI impurity of the node: 0.44
|       | Class distribution in the node: {0: 17, 1: 34}
|       | Predicted class: 1
|----- Split rule: Age > 29.5
|       | GINI impurity of the node: 0.32
|       | Class distribution in the node: {1: 61, 0: 15}
|       | Predicted class: 1
|----- Split rule: Age > 63.5
|       | GINI impurity of the node: 0.0
|       | Class distribution in the node: {0: 3}
|       | Predicted class: 0

```

Full decision tree; Snippet by author

The decision tree obtained from the scikit-learn implementation is identical:

```

|--- Fare <= 52.28
|   |--- Fare <= 10.48
|   |   |--- Age <= 32.50
|   |   |   |--- class: 0
|   |   |   |--- Age > 32.50
|   |   |   |--- class: 0
|   |   |--- Fare > 10.48
|   |   |   |--- Age <= 6.50
|   |   |   |   |--- class: 1
|   |   |   |   |--- Age > 6.50
|   |   |   |   |--- class: 0
|   |--- Fare > 52.28
|   |   |--- Age <= 63.50
|   |   |   |--- Age <= 29.50
|   |   |   |   |--- class: 1
|   |   |   |   |--- Age > 29.50
|   |   |   |   |--- class: 1
|   |   |--- Age > 63.50
|   |   |--- class: 0

```

Although, the scikit-learn implementation prints out less information than my implementation.

As it turns out, the best first initial split is the Fare feature at a value of 52.28 and not at the proposed Age feature at value 10.

The code that I have written builds the same trees as scikit-learn implementation and the predictions are the same. But the training time for the scikit-learn algorithm is much faster. But my goal was not to grow the trees faster. My goal was to write an understandable code for any machine learning enthusiast to have a better grasp of what is happening under the hood.

Feel free to create a pull request in this repo

<https://github.com/Eligijus112/decision-tree-python> if you see any bugs or just want to add functionalities.

Recursion

Python

Decision Tree

Machine Learning

Gini Impurity

Sign up for The Variable

By Towards Data Science

Every Thursday, the Variable delivers the very best of Towards Data Science: from hands-on tutorials and cutting-edge research to original features you don't want to miss. [Take a look.](#)

By signing up, you will create a Medium account if you don't already have one. Review our [Privacy Policy](#) for more information about our privacy practices.



Get this newsletter

[About](#) [Help](#) [Terms](#) [Privacy](#)

Get the Medium app

